

THE VANISHING CONSENSUS

PRO LEVEL LAB 02 • CASE NEX-071

Cross-Layer Trust-Chain Forensics: IoT Telemetry, Web3 Oracles, AI Model Poisoning, and Blockchain Consensus Divergence

Investigation Domain: WEB3 × AI × CYBERSECURITY × IoT

Case Identifier: CASE-NEX-071

Difficulty Level: Pro Certification (+300 XP)

Target Subject: Industrial IoT Gateway Ingestion Tampering, Oracle Aggregation Blindness, Neural Feedback Poisoning (MODEL-ORION), and BFT-PoS Validator Consensus Split

PREPARED BY:

Author: Lakshay Soni

Role: Security Analyst

Organization: TrinetLayer

Classification: RESTRICTED // SOC TIER 2 FORENSICS

Date of Incident & Report: October 2026

The Vanishing Consensus (Case NEX-071) investigates a sophisticated cross-layer trust-chain attack targeting industrial IoT infrastructure, Web3 oracle aggregation mechanisms, artificial intelligence sentinels, and Byzantine Fault Tolerant (BFT) blockchain consensus. At 03:12 UTC, 184 industrial sensors connected to GATEWAY-GW-184 began reporting perfectly synchronized telemetry (21.40°C, 412.00W, 0.00% jitter). This synthesized stream was ingested into NOVA-PRICE-ORACLE, where 4 independent oracle providers agreed on the corrupted data. Concurrently, the AI sentinel MODEL-ORION (v3.8.4) suppressed volatility alarms due to prior training dataset poisoning (EMB-IOT-9041). Consequently, validator nodes computed conflicting local state roots, causing a silent 3:2 consensus partition without triggering standard protocol crash errors.

TABLE OF CONTENTS

Section	Title	Domain Scope
1.0	Cover Page & Executive Dispatch	Administrative & Executive
2.0	Table of Contents	Navigation
3.0	Executive Summary & Incident Context	Strategic Cyber Overview
4.0	Lab Overview & Technical Parameters	Case NEX-071 Specifications
5.0	Complete Incident Briefing (The Core Anomaly)	Incident Telemetry
6.0	Complete Narrative Storyline & Characters	Story & Dialogue Records
7.0	Storyline & Forensic Investigation Timeline	Chronological Breakdown
8.0	Story → Technical Investigation Mapping Matrix	Cross-Domain Alignment
9.0	Chapter 1: The Signal That Lied	IoT Telemetry Forensics
10.0	Chapter 2: The Oracle That Saw Tomorrow	IoT × Web3 Oracle Security
11.0	Chapter 3: The Model That Learned the Attack	AI Model Poisoning Audit
12.0	Chapter 4: The Fork Nobody Saw	BFT Validator Consensus Split
13.0	Chapter 5: The Vanishing Consensus (Full Recovery)	Cross-Layer Containment
14.0	Complete 5-Chapter Progression Architecture	Technical Flowchart
15.0	Learning Objectives & Educational Competencies	Pedagogy & Training
16.0	Business & Operational Impact Analysis	Risk & Financial Exposure
17.0	Cross-Domain Attack Surface Analysis	IoT / Web3 / AI / Consensus
18.0	Threat Model & Adversary TTP Analysis	MITRE ATT&CK; Mapping
19.0	Cross-Domain Technical Architecture	System Schematic
20.0	IoT Architecture & Telemetry Ingestion Layer	Sensor & Gateway Telemetry
21.0	Web3 & Blockchain Architecture	Oracles, State Roots & BFT
22.0	AI Sentinel & Neural Decision Architecture	MODEL-ORION v3.8.4
23.0	Cybersecurity & Investigation Environment	SOC Desktop Suite
24.0	Practical Investigation Workflow	Step-by-Step Methodology
25.0	Forensic Evidence Vault & Cryptographic Artifacts	Tokens IOT-E11 to GOV-E15
26.0	End-to-End System Data Flow	Data Pipeline Diagram
27.0	Chapter Question & Forensic Validation Matrix	30 Verifiable Objectives
28.0	Final Multi-Layer Capstone Assessment	10-Question Evaluation
29.0	XP, Progress Tracking & Session Management	Scoring & Timer Mechanics
30.0	Achievements, Learning System & Leaderboard	User Motivation & Badges
31.0	Database Schema & Entity Relationship Model	SQLAlchemy Models
32.0	Codebase File Structure & Module Responsibilities	Repository Layout
33.0	Security Controls & Session Sandboxing	Authentication & CSRF
34.0	Responsive Multi-Platform Interface Design	Desktop / Tablet / Mobile
35.0	User Interface & Forensic Screenshot Documentation	SOC Tool Interfaces
36.0	Required Architectural & Flow Diagrams	19 Visual Schematics
37.0	Story + Technical Correlation Detailed Table	Pedagogical Integration
38.0	Complete Incident Response Timeline	Chronological Log
39.0	Final Case Resolution (The Master Flag)	NEXORA{v4n1sh1ng...}
40.0	Cybersecurity & Engineering Skills Developed	Hands-On Competencies
41.0	Technical, Operational & Strategic Lessons	Industry Takeaways
42.0	Current Limitations vs Proposed Enhancements	Gaps & Roadmap
43.0	Future Scope & Platform Evolution	Next-Gen Capabilities

44.0	Conclusion & Final Assessment	Closing Statement
45.0	Appendix A: Forensic File & Command Reference	Terminal & Log Reference
46.0	Appendix B: Evidence & Master Token Index	Cryptographic Hashes
47.0	Appendix C: Post-Investigation Debriefing Records	Incident Metrics

3.0 EXECUTIVE SUMMARY & INCIDENT OVERVIEW

The Vanishing Consensus (Case NEX-071) represents an advanced, multi-stage cybersecurity investigation that exposes the fragile trust boundaries linking physical Internet of Things (IoT) sensors, Web3 oracle data aggregators, artificial intelligence sentinels, and decentralized blockchain consensus protocols. Unlike conventional single-domain security scenarios—such as a smart contract reentrancy exploit or a standard credential theft attack—Case NEX-071 illustrates an adversary who deliberately refrains from attacking any single layer directly. Instead, the threat actor exploits the implicit trust relationships and data transformations between layers.

Why the Incident is Highly Suspicious: In physical sensor networks, real-world entropy, thermal variance, ambient vibrations, and local electrical load fluctuations naturally prevent identical telemetry readings across geographically dispersed nodes. In Case NEX-071, 184 industrial sensors across 14 global sites suddenly reported perfectly identical metrics (21.40°C temperature and 412.00W power draw) with exactly 0.00% measurement jitter. Downstream systems treated this unnatural synchronization not as an anomaly, but as verified reality, because the AI anomaly detection engine had been pre-conditioned via dataset poisoning to classify synthetic synchronization as benign.

4.0 LAB OVERVIEW & TECHNICAL PARAMETERS

Parameter	Specification Value	Investigation Significance
Lab Name	The Vanishing Consensus	Official Case Title
Case ID / Lab ID	Case NEX-071 / lab7	Database & Telemetry Identifier
Difficulty / Track	PRO / Professional Cybersecurity Track	Advanced Multi-Domain Lab
Total Experience (XP)	300 XP (+50 XP per verified task)	Player Progression Metric
Estimated Duration	55 - 65 Minutes (65m Session Auto-Restart)	Operational Time Budget
Investigation Scope	5 Chapters + 10-Question Multi-Layer Capstone	30 Core Tasks + Capstone Quiz
Technical Domains	IoT × Web3 × AI × Cybersecurity × Blockchain	Cross-Layer Converged Security
Target Systems	GATEWAY-GW-184, NOVA-PRICE-ORACLE, MODEL-ORION, BFT-POS	Primary Attacked Components
Available SOC Tools	Consensus Monitor, Burp Suite Proxy, Terminal, Case Files, Attack Graph	Integrated Desktop Workstation

5.0 COMPLETE INCIDENT BRIEFING // CASE NEX-071

■ CORE INCIDENT BRIEFING (AS PRESENTED IN SOC WORKSTATION)

03:12 AM. The blockchain monitoring system detects something that should be impossible.

A new block contains a transaction that appears to have happened before its own oracle price existed.

At almost the same time, several industrial IoT sensors connected to the Web3 infrastructure begin reporting perfectly synchronized telemetry.

Temperature. Power consumption. Location. Device health. Everything looks normal. Almost too normal.

Three validator nodes agree on the resulting blockchain state. Two reject it.

Meanwhile, the AI security engine reports: **'CONSENSUS HEALTHY — 98.7%'**

No validator appears compromised. No private key has been stolen. The IoT devices still appear online.

Yet the network is slowly beginning to disagree with itself. Someone isn't attacking the blockchain directly — someone is manipulating what the blockchain was told to believe.

6.0 COMPLETE STORYLINE & CHARACTER ROSTER

The narrative of **Case NEX-071** is structured around a multi-disciplinary security incident response team investigating a silent consensus divergence. Each team member brings deep domain-specific expertise, guiding the learner through physical IoT forensics, Web3 oracle data aggregation, AI model training feedback audits, and distributed consensus validation.

Character Name	Operational Role	Key Investigation Contribution	Primary Focus Area
Lakshay	Lead Incident Responder & Forensic Investigator	Directs cross-layer investigation, formulates hypotheses, correlates physical and on-chain evidence, executes recovery.	Cross-Domain Forensics & Recovery
Shivam	Distributed Systems & Web3 Infrastructure Lead	Analyzes IoT gateway logs, validator node state divergence, BFT-PoS consensus partition, and state root mismatches.	IoT Gateway & Validator Consensus
Mehak	AI Security & ML Governance Specialist	Discovers training dataset feedback poisoning (EMB-IOT-9041), audits neural inference confidence, and purges malicious weights.	AI Model Poisoning & Governance
Shanu	Threat Intelligence & Oracle Security Analyst	Investigates oracle feed dependencies, multi-provider aggregation blindness, and physical entropy absence.	Web3 Oracles & Threat Intel
MODEL-ORION	Autonomous AI Sentinel (v3.8.4)	Supervises live network transactions, suppresses circuit breakers due to poisoned baseline weights, reports false 98.7% health.	Simulated AI Sentinel Engine

7.0 STORYLINE & FORENSIC INVESTIGATION TIMELINE

Phase / Timestamp	Narrative Story Event	Technical Discovery / Clue	Investigative Action
03:10:44 UTC	Sensor Telemetry Ingestion	184 industrial sensors transmit valid local readings to GATEWAY-GW-184.	Hardware flash memory audit.
03:10:48 UTC	Gateway Telemetry Tampering	Gateway applies synthetic filter SYNTH_HARMONIC_V4, outputting 21.40°C & 412W.	Compare local flash vs gateway stream.
03:12:00 UTC	Oracle Ingestion & Broadcast	NOVA-PRICE-ORACLE consumes stream; 4 independent nodes agree on poisoned data.	Inspect oracle aggregation payload.
03:12:04 UTC	AI Sentinel Suppression	MODEL-ORION matches synthetic data to EMB-IOT-9041, reports NORMAL (98.7%).	Audit AI fine-tuning training corpus.
03:12:08 UTC	Validator Execution Divergence	3 nodes compute State Root 0x4f8e... (Accept); 2 nodes compute 0x98a2... (Reject).	Query 21-node validator state roots.
03:15:30 UTC	Cross-Layer Containment	GOV-NEX-071 proposal isolates gateway, purges model weights, restores consensus.	Execute governance proposal & verify flag.

9.0 CHAPTER 1: THE SIGNAL THAT LIED

Domain: IoT Security / Digital Forensics • Target Gateway: GATEWAY-GW-184 • Evidence Token: IOT-E11

Narrative Story & Dialogue Stream:

Shivam: 'The IoT gateway shows 184 active devices, but the telemetry patterns from several devices are nearly identical.'

Mehak: 'These sensors are in different locations. Their readings shouldn't synchronize this perfectly.'

Shanu: 'Perfect synchronization can be more suspicious than random noise. Real physical systems usually contain variation.'

Lakshay: 'Check whether the devices generated these readings locally or whether the gateway modified them.'

Shivam: 'That's the problem. The devices appear healthy, but the gateway logs show a different sequence of telemetry events.'

Mehak: 'Then the blockchain may not be the beginning of the attack. It may only be where the manipulated data became permanent.'

Chapter 1 Objectives & Verifiable Findings Matrix:

#	Objective Question	Verified Finding / Answer	Technical Explanation
1	Why are perfectly synchronized IoT readings suspicious?	Real physical systems contain variation	Natural physical environments exhibit entropy and measurement jitter.
2	Why is geographically identical telemetry unusual?	Different locations experience different environmental conditions	Dispersed sensors reflect differing ambient temperatures and power loads.
3	What should investigators compare first: device-generated telemetry or gateway telemetry?	Device-generated telemetry	Comparing local flash memory logs against gateway stream reveals in-transit tampering.
4	Why can a healthy-looking IoT device still be compromised?	Upstream gateway manipulates telemetry	Physical hardware remains intact while upstream gateway alters payload packets.
5	What evidence suggests that the gateway may have altered telemetry?	Gateway logs show a different sequence of events	Device memory logs showed normal variation while gateway recorded artificial stream.
6	If physical devices function normally but telemetry is manipulated, which layer failed?	Telemetry ingestion layer	The data ingestion and aggregation layer upstream of the blockchain was intercepted.

10.0 CHAPTER 2: THE ORACLE THAT SAW TOMORROW

Domain: IoT × Web3 Oracle Security • Target Oracle: NOVA-PRICE-ORACLE • Evidence Token: ORACLE-E12

Narrative Story & Dialogue Stream:

Shanu: 'The oracle isn't receiving raw blockchain data. It's consuming external telemetry.'

Shivam: 'And that telemetry originates from the IoT gateway.'

Mehak: 'Then the oracle providers aren't independent if they're all consuming the same manipulated telemetry stream.'

Lakshay: 'Exactly. The attack moved from the physical world into the Web3 trust layer.'

Shanu: 'Four providers can appear to agree while actually repeating the same poisoned information.'

Mehak: 'Which means the blockchain may be reaching consensus on data that was already compromised before it entered the chain.'

Chapter 2 Objectives & Verifiable Findings Matrix:

#	Objective Question	Verified Finding / Answer	Technical Explanation
1	Where does NOVA-PRICE-ORACLE obtain its external data?	External telemetry aggregation layer	Aggregates external telemetry streams from industrial IoT gateways.
2	Why does compromised IoT telemetry affect Web3 systems?	Oracles consume external telemetry as truth	Smart contracts cannot fetch off-chain data independently and trust oracle feeds.

3	Why can multiple oracle providers still represent one single source?	They consume the same upstream IoT telemetry stream	Multiple independent nodes consuming a single corrupted stream replicate identical error.
4	What happens when manipulated telemetry is aggregated before reaching chain?	Poisoned aggregate enters oracle as trusted state	Pre-chain math averages poisoned telemetry, cementing falsified value on-chain.
5	Why is 'multiple providers agree' insufficient for data integrity?	Providers repeat the same corrupted upstream source	Node consensus only proves agreement on received data, not external physical truth.
6	At what point did the attack cross from IoT into Web3 security?	When compromised IoT telemetry became trusted external oracle data	The boundary transition occurred when the oracle feed accepted poisoned gateway data.

11.0 CHAPTER 3: THE MODEL THAT LEARNED THE ATTACK

Domain: IoT AI Security / Model Poisoning • AI Sentinel: MODEL-ORION (v3.8.4) • Evidence Token: AI-E13

Narrative Story & Dialogue Stream:

Shanu: 'ORION saw the synchronized IoT telemetry, but it classified the pattern as normal.'

Lakshay: 'Can the model distinguish physical anomalies from fabricated telemetry?'

Shanu: 'Only if its training data taught it what both look like.'

Mehak: 'I found repeated synthetic device events in the training feedback.'

Shivam: 'So someone wasn't only manipulating live IoT telemetry. They were preparing the AI to trust similar behavior.'

Lakshay: 'The physical layer was poisoned first. Then the AI was trained not to question the poison.'

Chapter 3 Objectives & Verifiable Findings Matrix:

#	Objective Question	Verified Finding / Answer	Technical Explanation
1	What IoT behavior did ORION classify as normal?	Synchronized IoT telemetry	MODEL-ORION classified synchronized readings as NORMAL_NETWORK_VARIANCE.
2	Why is synchronized telemetry useful for detecting IoT manipulation?	It indicates synthetic data injection	In genuine physical networks, cross-sensor synchronization indicates replay attacks.
3	How can poisoned training data affect IoT security?	AI classifies malicious patterns as normal	When datasets contain adversarial samples labelled benign, AI ignores live attacks.
4	What is the relationship between IoT telemetry and the AI anomaly detector?	IoT telemetry serves as feature inputs for AI	Incoming sensor streams form primary input feature vectors for neural evaluation.
5	Why can a high-confidence AI result still be wrong?	Model was trained on poisoned data	Statistical confidence only reflects alignment with weights, which may be corrupted.
6	If an AI is trained on manipulated telemetry, can it detect it later?	No, because it learned manipulated patterns as legitimate	The model internalizes malicious patterns as baseline, blinding it during live inference.

12.0 CHAPTER 4: THE FORK NOBODY SAW

Domain: IoT × AI × Blockchain Consensus • Topology: 21-Node BFT Grid • Evidence Token: CONSENSUS-E14

Narrative Story & Dialogue Stream:

Shivam: 'The validators are no longer processing exactly the same derived state.'

Mehak: 'Because the upstream telemetry was inconsistent?'

Shivam: 'Exactly. Different data paths produced slightly different oracle states.'

Shanu: 'And ORION isn't escalating the difference because the model considers the telemetry pattern normal.'

Lakshay: 'So IoT manipulation created the input anomaly, AI suppressed the warning, and the blockchain inherited the disagreement.'

Mehak: 'Three separate security layers failed in sequence.'

Chapter 4 Objectives & Verifiable Findings Matrix:

#	Objective Question	Verified Finding / Answer	Technical Explanation
1	How did IoT telemetry ultimately influence validator state?	Altering smart contract inputs through poisoned oracle	Manipulated telemetry fed oracle, injecting false price state into contract execution.
2	Why did different validators receive different derived states?	Conflicting oracle data paths and timing across nodes	Network propagation latency caused some nodes to compute Root A while others Root B.
3	How did AI failure contribute to the consensus problem?	ORION suppressed anomaly warnings by classifying patterns as normal	Because AI classified telemetry as normal, automated circuit breakers never paused blocks.

4	Why didn't the blockchain immediately detect the issue?	Validators executed deterministically from local inputs without protocol errors	Each node executed valid VM bytecode; error was in truthfulness of inputs, not syntax.
5	Which security layers were involved in the attack?	IoT, Oracle Web3, AI, Blockchain	The attack chained across physical sensors, Web3 oracles, AI sentinels, and BFT consensus.
6	Why is cross-layer attack harder to detect than single-layer?	Each layer appears individually healthy while trust boundaries are exploited	No component crashes; exploit occurs across inter-layer trust assumptions.

13.0 CHAPTER 5: THE VANISHING CONSENSUS (FULL ATTACK RECONSTRUCTION)

Domain: Full Cross-Layer Attack Reconstruction • Proposal: GOV-NEX-071 • Evidence Token: GOV-E15

Narrative Story & Dialogue Stream:

Lakshay: 'The attack started outside the blockchain.'

Shivam: 'At the IoT telemetry layer.'

Mehak: 'The manipulated data then entered the oracle ecosystem.'

Shanu: 'And the poisoned AI model helped hide the anomaly.'

Shivam: 'Different validators eventually processed different states.'

Lakshay: 'So the blockchain wasn't directly hacked.'

Mehak: 'It was manipulated through the trust chain connecting IoT, AI, Web3 and blockchain.'

Chapter 5 Objectives & Verifiable Findings Matrix:

#	Objective Question	Verified Finding / Answer	Technical Explanation
1	Where did the attack actually begin?	IoT telemetry layer	The root origin was the compromised telemetry stream on GATEWAY-GW-184.
2	Why was the IoT layer critical to the attack?	It provided the initial poisoned telemetry feeding downstream systems	Without manipulated sensor stream, downstream oracle and AI produce no fraud state.
3	How did manipulated telemetry reach the blockchain?	Via Web3 oracle aggregation feed	The oracle pipeline ingested gateway data and broadcasted it as verified on-chain state.
4	How did AI help hide the attack?	By classifying anomalous telemetry and spikes as legitimate	Poisoned neural network categorized anomaly as benign, preventing SOC alarms.
5	Why did validator disagreement occur?	Validators processed conflicting oracle-derived states	Differences in oracle propagation timing produced divergent state transitions.
6	Why is the attack best described as a 'cross-layer trust-chain attack'?	Exploited trust dependencies connecting IoT, Web3, AI, and Blockchain	Chained vulnerabilities across IoT integrity, oracle trust, AI alignment, and BFT consensus.

28.0 FINAL MULTI-LAYER CAPSTONE ASSESSMENT (10 QUESTIONS)

Evaluation Standard: 10 Multi-Choice Forensic Questions • Passing Threshold: ≥ 70% • XP Reward: +300 XP

#	Capstone Evaluation Question	Correct Verified Option	Evaluation Rationale
1	What was the earliest suspicious indicator in Case NEX-071?	B) Perfectly synchronized IoT telemetry across sensors	Physical environments have natural entropy; identical readings indicate synthesis.
2	Why was the IoT telemetry critical to the attack?	A) It became manipulated data source feeding Web3/blockchain	Physical sensor stream is the ground-truth feed for smart contract oracles.
3	Why did oracle provider agreement fail to guarantee truth?	A) Multiple providers depended on same compromised source	Agreement only proves replication consistency, not source physical integrity.
4	How was MODEL-ORION compromised?	A) Training feedback contained synthetic IoT data labelled legitimate	Adversary poisoned fine-tuning dataset with EMB-IOT-9041 signature.
5	Why could ORION confidently classify attack as normal?	A) Model learned manipulated patterns as normal baseline	High statistical confidence reflected alignment with poisoned training weights.
6	How did IoT manipulation affect blockchain consensus?	A) Altered oracle states causing validators to process different states	Inconsistent oracle timing produced divergent local EVM state roots.
7	What made the attack difficult to detect?	A) Each layer appeared healthy while trust boundaries were exploited	No single module crashed; trust assumptions between layers were breached.

8	Which sequence best represents the attack?	B) IoT → Oracle → AI → Validator divergence	Physical telemetry poisoned → Oracle ingested → AI suppressed → Chain split.
9	Did the attacker need to directly control or hack the blockchain?	B) No — manipulated trusted information upstream of consensus	Consensus algorithms executed properly on corrupted input data.
10	What is the deepest security lesson from Case NEX-071?	D) Security depends on protecting trust boundaries across IoT, AI, Web3, Blockchain	End-to-end security requires cross-layer verification and physical entropy audits.

17.0 CROSS-DOMAIN ATTACK SURFACE & 18.0 THREAT MODEL

Cross-Domain Attack Surface Analysis:

- **Physical & IoT Ingestion Surface:** Industrial telemetry gateway (GATEWAY-GW-184) lacks hardware-enforced measurement entropy validation, allowing synthetic payload overrides (SYNTH_HARMONIC_V4).
- **Web3 Oracle Aggregation Surface:** NOVA-PRICE-ORACLE relies on multi-node multi-signature consensus over an unverified single upstream telemetry aggregator, confounding consensus with ground truth.
- **AI Sentinel Decision Surface:** MODEL-ORION consumes automated training feedback loops without adversarial data sanitation, allowing backdoor baseline poisoning (EMB-IOT-9041).
- **Blockchain Consensus Surface:** Byzantine Fault Tolerant (BFT) Proof-of-Stake state machine deterministically processes external oracle state transitions, propagating input skew into state root divergence (3:2 split).

Threat Model & Adversary TTP Matrix (MITRE ATT&CK; Alignment):

Dimension	Adversary Realization in Case NEX-071	Mitigation & Containment Strategy
Threat Actor	Advanced Persistent Threat (APT) targeting Distributed Trust	Cross-layer forensic telemetry & multi-domain threat correlation.
Initial Entry	Compromised Ingestion Daemon on GATEWAY-GW-184	Hardware root-of-trust (TPM) & signed local flash logs.
Attack Vector	Synthetic Sensor Synchrony & Upstream Stream Overwrite	Physical entropy and environmental jitter verification filters.
AI Exploitation	Dataset Poisoning via EMB-IOT-9041 (1,200 synthetic vectors)	Adversarial data sanitation & offline human-in-the-loop retraining.
Web3 Exploitation	Oracle Aggregator Blindness (4/4 False Agreement)	Multi-source physical diversity & outlier rejection algorithms.
Consensus Impact	3:2 Validator State Root Divergence (Block #982741)	Automated EVM circuit breakers & emergency governance rollback.

19.0 CROSS-DOMAIN TECHNICAL ARCHITECTURE SCHEMATIC

[illegible]

23.0 INVESTIGATION ENVIRONMENT & SOC DESKTOP TOOLS

SOC Tool Name	Interface & Tabs	Investigative Purpose in Case NEX-071	Active Chapter(s)
Consensus Monitor	Browser Window (6 Tabs: IoT, Oracle, AI, Consensus, Governance, DevTools)	Inspect live sensor telemetry, oracle prices, AI confidence, validator grid status, and governance proposals.	Chapters 1, 2, 3, 4, 5
Burp Suite Proxy	Burp Window (Proxy HTTP History, Repeater, Inspector, Target Scope)	Capture, inspect, and repeat raw HTTP JSON payloads between IoT gateways, oracles, AI sentinels, and validators.	Chapters 1, 2, 3, 4, 5
Forensic Terminal	bash Terminal (`investigator@workstation:~\$`)	Execute specialized commands: `iot-telemetry`, `oracle-status`, `ai-audit`, `validator-status`, `python inspect_attack_chain.py`.	Chapters 1, 2, 3, 4, 5
Case Files Explorer	File Manager (`/var/log/consensus`)	Review raw logs: `iot-telemetry-gw184.log`, `oracle-aggregation.json`, `orion-training-poison.log`, `validator-divergence.log`.	Chapters 1, 2, 3, 4, 5
Attack Graph	Incident Flowchart Modal	Visualize cross-layer attack progression from physical IoT to oracle, AI, consensus split, and governance containment.	Chapters 1, 2, 3, 4, 5
Evidence Vault	Cryptographic Artifact Locker	Secure SHA256-verified tokens: `IOT-E11`, `ORACLE-E12`, `AI-E13`, `CONSENSUS-E14`, `GOV-E15`.	Chapters 1, 2, 3, 4, 5
Notes Scratchpad	Investigator Notepad	Persist analyst observations, sensor IDs, oracle metrics, and hypotheses in browser localStorage.	Chapters 1, 2, 3, 4, 5

25.0 FORENSIC EVIDENCE VAULT & CRYPTOGRAPHIC ARTIFACTS

Evidence Token	Evidence Name	Cryptographic SHA256 Hash	Forensic Description & Source
IOT-E11	Synchronized IoT Telemetry Log	0x7a8b1102e4d91c28f731	Capture from GATEWAY-GW-184 showing artificial synchronization across 184 dispersed sensors.
ORACLE-E12	Aggregated Oracle Data Feed	0x9c3d5412a819bb204f62	Payload from NOVA-PRICE-ORACLE showing how corrupted IoT telemetry became on-chain state.
AI-E13	Poisoned AI Training Feedback	0x1f4a8831ef78c90231aa	Training audit log for MODEL-ORION revealing synthetic device events (EMB-IOT-9041).
CONSENSUS-E14	Validator State Root Trace	0x3e2b6904d9a17730bc12	Consensus telemetry showing 3:2 validator partition (VALIDATOR-V03 vs VALIDATOR-V05).
GOV-E15	Unified Containment Proposal	0x8d1e7752fc304192b091	Emergency governance execution (GOV-NEX-071) restoring unified consensus across all 21 nodes.

31.0 DATABASE SCHEMA & 32.0 REPOSITORY CODE STRUCTURE

File / Module Path	Primary Responsibility in Lab 2	Key Functions & Implementations
seed_vanishing_consensus.py	Database seeder for Case NEX-071	Seeds `lab7`, 5 Missions, 30 MissionQuizzes, 5 Hints, 5 Evidence items, Master Flag.
templates/vanishing_consensus_workstation.html	Core investigation SOC workstation UI	Renders two-column responsive layout, 5 accordion tasks, 6 browser tabs, Burp, Terminal, Files.
templates/vanishing_consensus_post_investigation.html	Post-investigation debrief & celebration	Renders completion celebration, 7-stage attack chain reconstruction, score and rank.
static/js/vanishing_consensus_desktop.js	Client-side interactive forensic desktop OS	Manages 65m timer, window z-index, terminal commands, Burp repeater, capstone validation.
static/css/pages/vanishing_consensus_desktop.css	Cyberpunk / dark-theme styling	Defines responsive layouts, telemetry cards, diff tables, node grid, Burp theme.
routes/labs.py	Flask HTTP routes & API handlers	Handles `/lab/vanishing-consensus/workstation`, evaluation APIs, and post-investigation.

39.0 FINAL CASE RESOLUTION & MASTER INVESTIGATION FLAG

■ CASE NEX-071 MASTER FLAG & GOVERNANCE RECOVERY

Investigation Conclusion (FINAL REVELATION):

The blockchain was never directly broken. The AI was never simply switched off. The IoT devices were never physically destroyed. The adversary compromised something far more critical: **THE TRUST RELATIONSHIPS CONNECTING THE LAYERS.**

By isolating GATEWAY-GW-184, purging poisoned baseline weights (EMB-IOT-9041) from MODEL-ORION, adding physical entropy validation to NOVA-PRICE-ORACLE, and replaying state to Block #982740 via governance proposal GOV-NEX-071, unified consensus was successfully restored across all 21 validator nodes.

MASTER CASE FLAG: NEXORA{v4n1sh1ng_c0ns3nsus_n3x071}

44.0 CONCLUSION & ANALYST SIGN-OFF

PRO LEVEL LAB 02: THE VANISHING CONSENSUS (Case NEX-071) establishes a benchmark for next-generation cybersecurity education. By moving beyond isolated vulnerabilities and confronting learners with a realistic, multi-layered trust-chain failure across IoT, Web3, AI, and Blockchain consensus, the lab instills critical forensic intuition: *consensus is not synonymous with truth, and system security depends on defending inter-layer trust boundaries.*

Report Prepared By	Technical Role	Organization	Case File Status
Lakshay Soni	Security Analyst	TrinetLayer	OFFICIALLY RESOLVED // FLAG VERIFIED